

MultipleLinearRegression

December 5, 2023

0.1 House Price Prediction

0.1.1 Import

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
```

```
[2]: df = pd.read_csv('housing_price_dataset.csv')
df.head(5)
```

```
[2]:
```

	SquareFeet	Bedrooms	Bathrooms	Neighborhood	YearBuilt	Price
0	2126	4	1	Rural	1969	215355.283618
1	2459	3	2	Rural	1980	195014.221626
2	1860	2	1	Suburb	1970	306891.012076
3	2294	2	1	Urban	1996	206786.787153
4	2130	5	2	Suburb	2001	272436.239065

0.1.2 Data Cleaning

```
[3]: df.shape
```

```
[3]: (50000, 6)
```

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   SquareFeet      50000 non-null  int64
1   Bedrooms        50000 non-null  int64
2   Bathrooms       50000 non-null  int64
```

```

3   Neighborhood  50000 non-null  object
4   YearBuilt     50000 non-null  int64
5   Price         50000 non-null  float64
dtypes: float64(1), int64(4), object(1)
memory usage: 2.3+ MB

```

```
[5]: df.describe()
```

```

[5]:      SquareFeet    Bedrooms    Bathrooms    YearBuilt      Price
count  50000.000000  50000.000000  50000.000000  50000.000000  50000.000000
mean    2006.374680    3.498700    1.995420    1985.404420  224827.325151
std      575.513241    1.116326    0.815851    20.719377   76141.842966
min     1000.000000    2.000000    1.000000    1950.000000 -36588.165397
25%     1513.000000    3.000000    1.000000    1967.000000  169955.860225
50%     2007.000000    3.000000    2.000000    1985.000000  225052.141166
75%     2506.000000    4.000000    3.000000    2003.000000  279373.630052
max     2999.000000    5.000000    3.000000    2021.000000  492195.259972

```

```
[6]: df.dropna()
```

```

[6]:      SquareFeet  Bedrooms  Bathrooms  Neighborhood  YearBuilt      Price
0           2126          4           1          Rural      1969  215355.283618
1           2459          3           2          Rural      1980  195014.221626
2           1860          2           1          Suburb      1970  306891.012076
3           2294          2           1          Urban      1996  206786.787153
4           2130          5           2          Suburb      2001  272436.239065
...          ...          ...          ...          ...          ...
49995         1282          5           3          Rural      1975  100080.865895
49996         2854          2           2          Suburb      1988  374507.656727
49997         2979          5           3          Suburb      1962  384110.555590
49998         2596          5           2          Rural      1984  380512.685957
49999         1572          5           3          Rural      2011  221618.583218

```

```
[50000 rows x 6 columns]
```

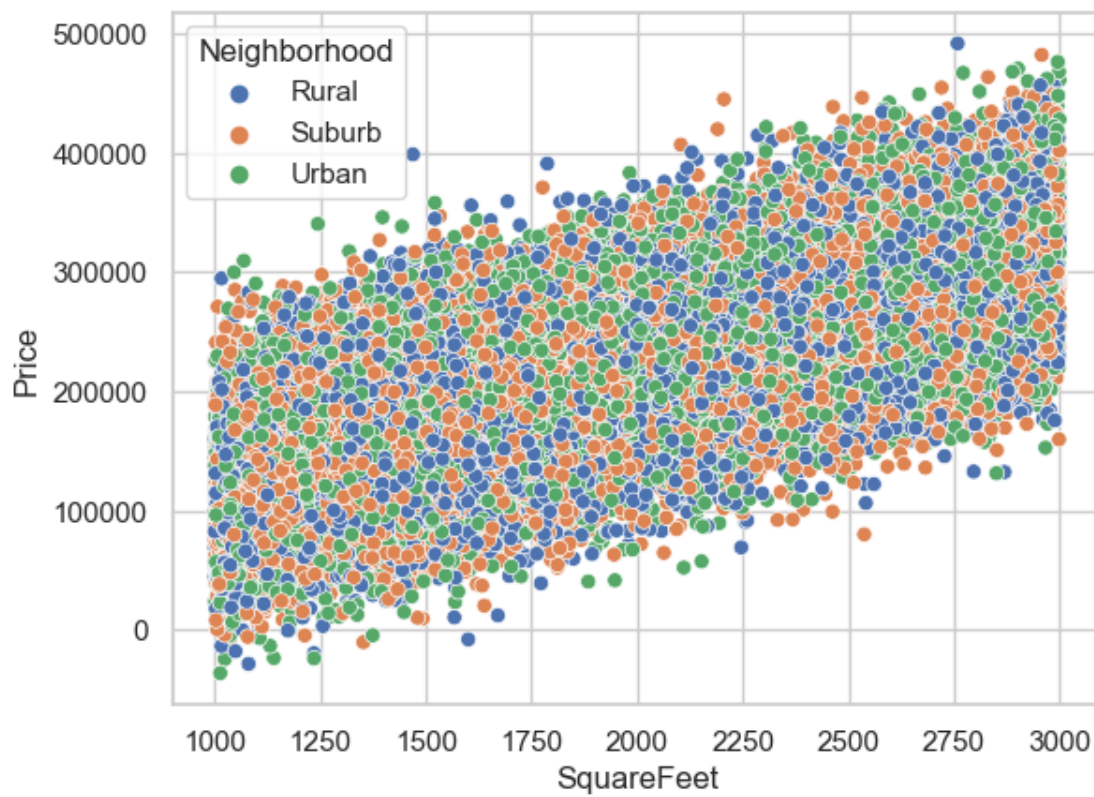
0.1.3 Visualization

```

[7]: sns.set(style = 'whitegrid')
      sns.scatterplot(x = 'SquareFeet', y = 'Price', data = df, hue = 'Neighborhood')

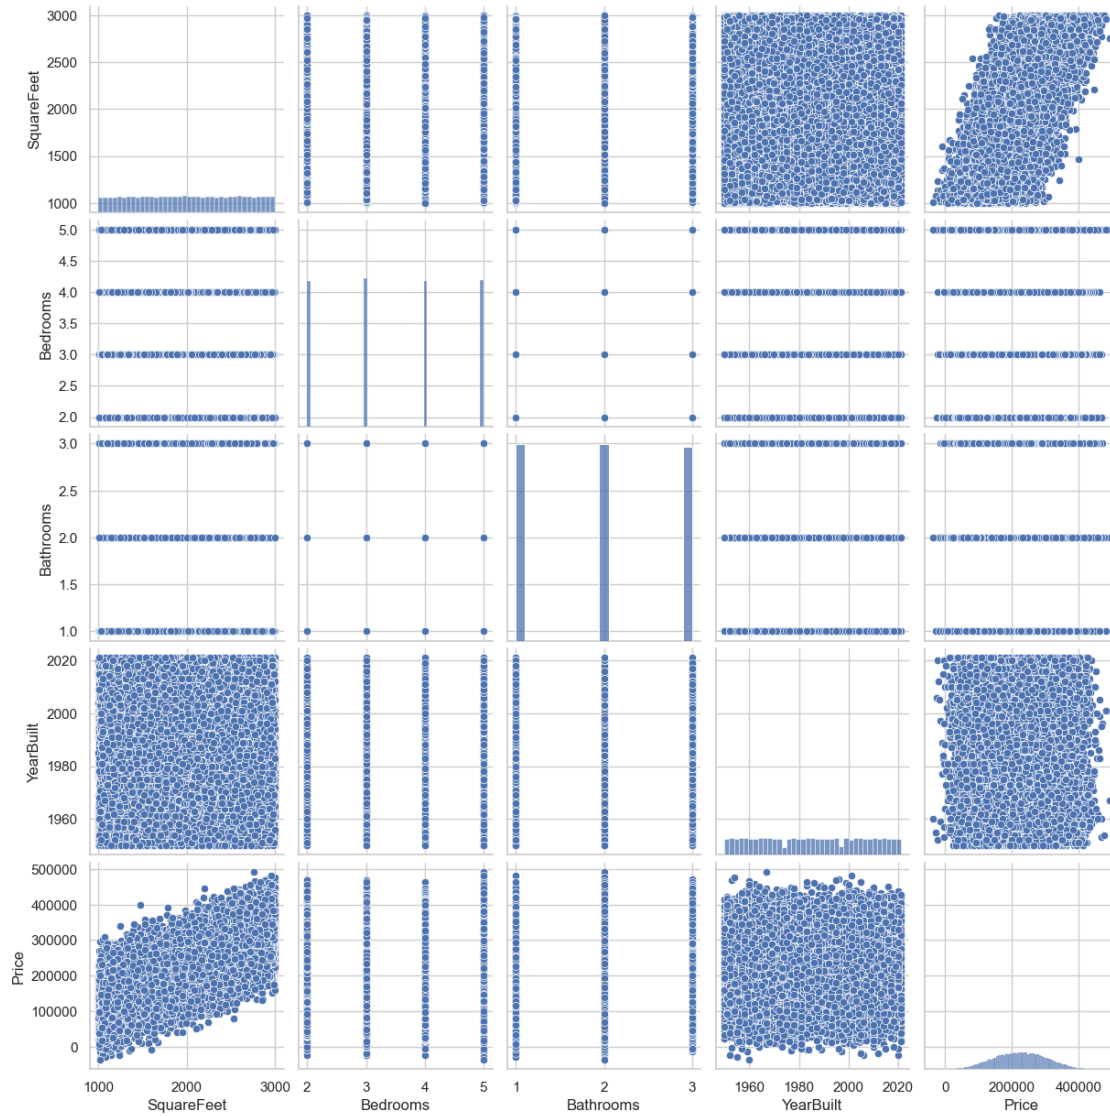
```

```
[7]: <Axes: xlabel='SquareFeet', ylabel='Price'>
```




```
[8]: sns.pairplot(df)
```

```
[8]: <seaborn.axisgrid.PairGrid at 0x219db262150>
```



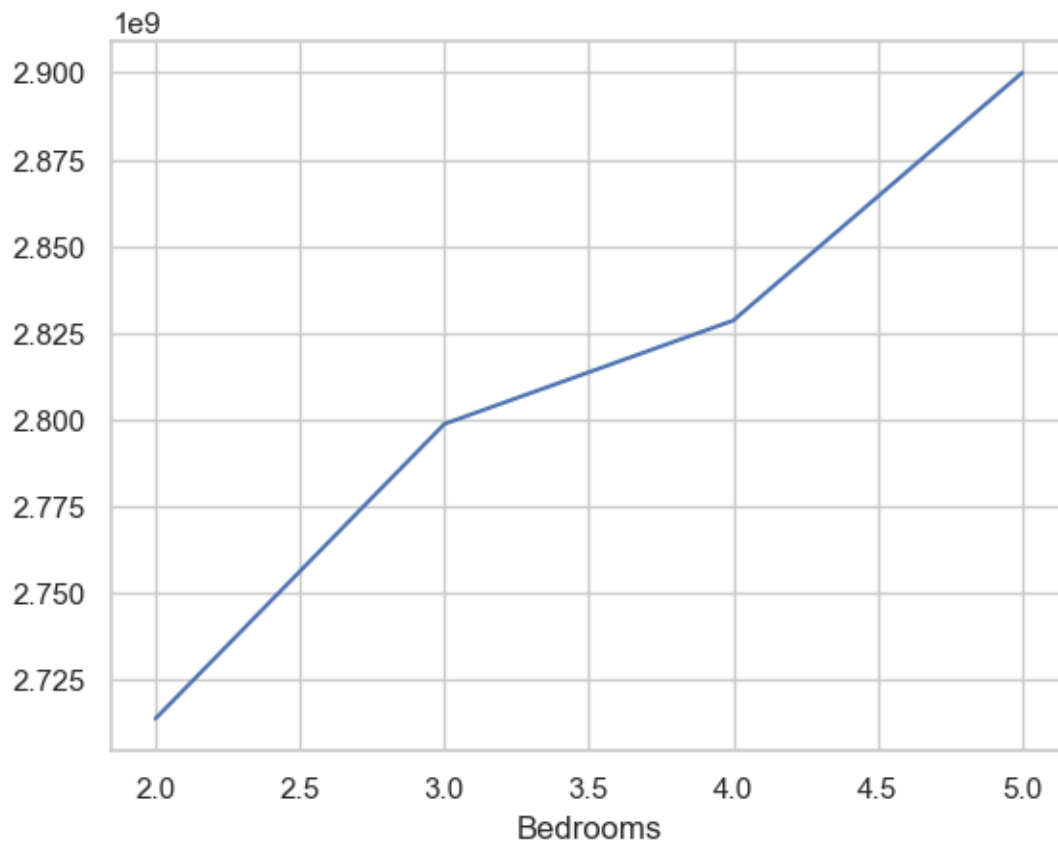
```
[9]: df.groupby('YearBuilt')['Price'].sum().plot(kind = 'line')
```

```
[9]: <Axes: xlabel='YearBuilt'>
```




```
[10]: df.groupby('Bedrooms')['Price'].sum().plot(kind = 'line')
```

```
[10]: <Axes: xlabel='Bedrooms'>
```



```
[11]: df.groupby('Bathrooms')['Price'].sum().plot(kind = 'line')
```

```
[11]: <Axes: xlabel='Bathrooms'>
```



One-hot Encoding

```
[12]: # Since the Neighbourhoods are of type object, convert that to numeric type..
df = pd.get_dummies(df,columns=['Neighborhood'],drop_first=True, dtype=int)
df.head()
```

```
[12]:
```

	SquareFeet	Bedrooms	Bathrooms	YearBuilt	Price \
0	2126	4	1	1969	215355.283618
1	2459	3	2	1980	195014.221626
2	1860	2	1	1970	306891.012076
3	2294	2	1	1996	206786.787153
4	2130	5	2	2001	272436.239065

	Neighborhood_Suburb	Neighborhood_Urban
0	0	0
1	0	0
2	1	0
3	0	1
4	1	0

```
[13]: Y = df['Price']
      X = df.drop('Price',axis=1)
```

```
[14]: X.columns
```

```
[14]: Index(['SquareFeet', 'Bedrooms', 'Bathrooms', 'YearBuilt',
          'Neighborhood_Suburb', 'Neighborhood_Urban'],
          dtype='object')
```

Scaling

```
[15]: ss = StandardScaler()
      X[['SquareFeet', 'Bedrooms', 'Bathrooms','YearBuilt']] = ss.
      ↪fit_transform(X[['SquareFeet', 'Bedrooms', 'Bathrooms','YearBuilt']])
      X.head()
```

```
[15]:   SquareFeet  Bedrooms  Bathrooms  YearBuilt  Neighborhood_Suburb  \
0    0.207861  0.449067  -1.220113  -0.791751                0
1    0.786480 -0.446738   0.005614  -0.260842                0
2   -0.254340 -1.342543  -1.220113  -0.743486                1
3    0.499777 -1.342543  -1.220113   0.511390                0
4    0.214811  1.344872   0.005614   0.752713                1

      Neighborhood_Urban
0                      0
1                      0
2                      0
3                      1
4                      0
```

0.1.4 Model Training

```
[16]: X_train, X_test, Y_train, Y_test = train_test_split(X,Y)
```

```
[17]: X_train.shape,X_test.shape
```

```
[17]: ((37500, 6), (12500, 6))
```

```
[18]: lr = LinearRegression()
      lr.fit(X_train,Y_train)
```

```
[18]: LinearRegression()
```

```
[24]: Y_pred_train = lr.predict(X_train)
      Y_pred_test = lr.predict(X_test)
      print(f'Training RMSE: {np.sqrt(mean_squared_error(Y_pred_train,Y_train))}')
      print(f'Test RMSE: {np.sqrt(mean_squared_error(Y_pred_train,Y_train))}')
```


Training RMSE: 49871.25195609414

Test RMSE: 49871.25195609414

- Prediction Analysis

```
[22]: sns.kdeplot(Y_train)
sns.kdeplot(Y_pred_train)
plt.title('Training')
plt.show()
sns.kdeplot(Y_test)
sns.kdeplot(Y_pred_test)
plt.title('Test')
plt.show()
```

